


Model Selection:

- Recall ERM:

$$\min_{\vec{w}} \frac{1}{n} \sum_{i=1}^n \underbrace{\ell(h_{\vec{w}}(x_i), y_i)}_{\text{Loss}} + \underbrace{\lambda r(w)}_{\text{Regularizer}}$$

- Under fitting: The solution is too simple.  
The training and test error will be high.
- Overfitting: The solution is too complex.  
The training error decrease over time, the test error will begin to increase.



- Identifying sweet spot: Divide data into training and validation portions. Train on the "training" split and evaluate it on the "validation" split, for various values of  $\lambda$  ( $10^{-5}, 10^{-4}, 10^{-2}, 10^{-1}, 10^0, 10^1, 10^2$ )



← We may overfit the validation set.

## K-fold cross validation:


- Divide your training data into  $k$  parts.
- Train on  $k-1$  parts and leave one out as validation set. Do this  $k$  times and average the validation error across all runs (ideally,  $k=n$ ).

## Telescopic search for $\lambda$ :

- Do two search:
  - 1st step: find the best order of magnitude for  $\lambda$ 
    - For example, first we try 0.01, 0.1, 1, 10, 100.
 

the best  
↓
  - 2nd step: do a more fine-grained search around the best  $\lambda$  so far
    - Then, we try 5, 10, 15, 20, 25, ..., 95 to test values around 10.

Early Stopping: Stop our optimization after  $M \geq 0$  number of gradient steps, even if optimization has not converged yet.



## Kernels:


See visualization at: <https://www.youtube.com/watch?v=3liCbRZPrZA>

- A way to incorporate non-linearities into most linear classifiers.

Feature Transformation:

$$\vec{x} \rightarrow \phi(\vec{x})$$

$$\begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \rightarrow \begin{pmatrix} x_1 \\ x_2 \\ x_1^2 + x_2^2 \end{pmatrix}$$



Extreme case:

$$\begin{pmatrix} x_1 \\ \vdots \\ x_d \end{pmatrix} \rightarrow \begin{pmatrix} 1 \\ x_1 \\ \vdots \\ x_d \\ x_1 x_2 \\ \vdots \\ x_{d-1} x_d \\ x_1 x_d \dots x_d \end{pmatrix} \in 2^d$$